



USER SESSION RECORDING

An Open Source solution

Fraser Tweedale

@hackuador

2017-10-22

ABOUT ME

- Working at Red Hat, Platform Engineering (Security)
- FreeIPA and Dogtag Certificate System

WHY?

THERE IS A DEMAND

Customers need to...

- comply with government or industry regulations
- track what contractors do on our systems
- know who broke our server, and how

AND A DREAM

What companies and governments want:

- Record everything users do
- Store that somewhere safe
- Let us find who did *that thing*
- Show us how they did it

THERE IS A SUPPLY

A number of commercial offerings:

- From application-level proxies on dedicated hardware
- To user-space processes on the target system
- Recording keystrokes, display, commands, apps, URLs, etc.
- Integrated with identity management, and access control
- With central storage, searching, and playback

BUT NOT GOOD ENOUGH

Customers are not satisfied:

- Expensive
- Can't fix it yourself
- Can't improve it yourself

WHAT CAN BE BETTER?

The customers want:

- Lower costs
- Open Source, so they can fix, or at least understand it better
- Commercial support

WAIT, WE HAVE IT ALREADY!

Nope, not really:

- `script(1)` plus duct tape
 - popular, but not security-oriented; lots of DIY
- `sudo(8)` I/O logging
 - security-oriented, has searching, but not centralised
- TTY audit with `auditd(8)`
 - security-oriented, can be centralized, only records input

SO, WHAT DO WE NEED?

1. Record terminal I/O from userspace
2. Prompt, secure, centralised logging
3. Log kernel audit events, too
4. Search & playback of recorded sessions; correlation
5. Centralised control

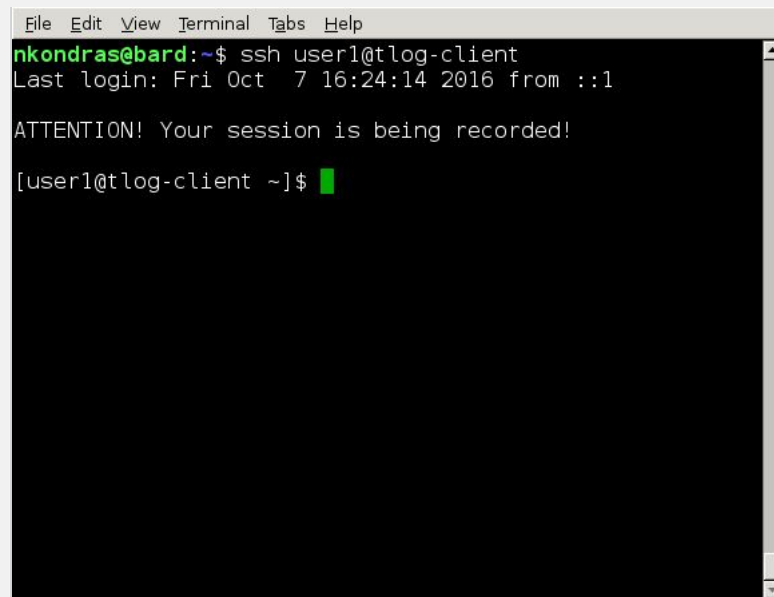
1. RECORDING TERMINAL I/O

1. RECORD SESSION I/O

We made a tool for that - **tlog**

<http://scribery.github.io/tlog>

- A shim between the terminal and the shell, started at login
- Converts terminal activity to JSON
- Log to **file**, **syslog** or **journal**
- Playback to terminal



```
File Edit View Terminal Tabs Help
nkondras@bard:~$ ssh user1@tlog-client
Last login: Fri Oct  7 16:24:14 2016 from ::1

ATTENTION! Your session is being recorded!

[user1@tlog-client ~]$
```

EXTRA TLOG FEATURES


Also control:

- What to record: input/output/window resizes
- “*You are being recorded*” notice
- Low latency vs. low overhead

Basic playback on the terminal:

- From elasticsearch, journal or file

TLOG MINIMAL SETUP




HOW TLOG WORKS?

Console login example

Starting a console session:

1. User authenticates to **login** via **PAM**
2. **NSS** tells **login**: **tlog** is the shell
3. **login** starts **tlog**
4. env/config tell **tlog** the actual shell
5. **tlog** starts the actual shell in a **pty**
6. **tlog** logs everything passing between its **terminal** and the **pty**, via **syslog(3)** or **sd-journal(3)**



TLOG SCHEMA

Optimized for streaming and searching:

- Chopped into messages for streaming, which can be merged
- Input and output stored separately
- All I/O preserved
- Invalid UTF-8 stored separately
- Timing separate, ms precision
- Window resizes preserved

```
{
  "ver"      : "2.2",
  "host"     : "tlog-client.example.com",
  "rec"      : "c8aa248c81264f5d98d1...",
  "user"     : "user1",
  "term"     : "xterm",
  "session"  : 23,
  "id"       : 1,
  "pos"      : 0,
  "timing"    : "=56x22+98>23",
  "in_txt"   : "",
  "in_bin"   : [ ],
  "out_txt"  : "[user1@tlog-client ~]$ ",
  "out_bin"  : [ ]
}
```


JOURNAL FORMAT

Exposes key fields

```
{
  "ver"      : "2.2",
  "host"     : "tlog-client.example.com",
  "rec"      : "c8aa248c81264f5d98d1...",
  "user"     : "user1",
  "term"     : "xterm",
  "session"  : 23,
  "id"       : 1,
  "pos"      : 0,
  "timing"   : "=56x22+98>23",
  "in_txt"   : "",
  "in_bin"   : [ ],
  "out_txt"  : "[user1@tlog-client ~]$ ",
  "out_bin"  : [ ]
}
```

```
_AUDIT_SESSION=23
_AUDIT_LOGINUID=1000
TLOG_REC=c8aa248c81264f5d98d1...
TLOG_USER=user1
_UID=987
TLOG_SESSION=23
TLOG_ID=1
_COMM=tlog-rec-sess
_EXE=/usr/bin/tlog-rec-session
_TRANSPORT=journal
MESSAGE={"ver":"2.2","host":...}
SYSLOG_IDENTIFIER=tlog-rec-session
...
```

2. LOGGING INFRASTRUCTURE

2. LOGGING INFRASTRUCTURE

What to take out of the store/search/analyze zoo?

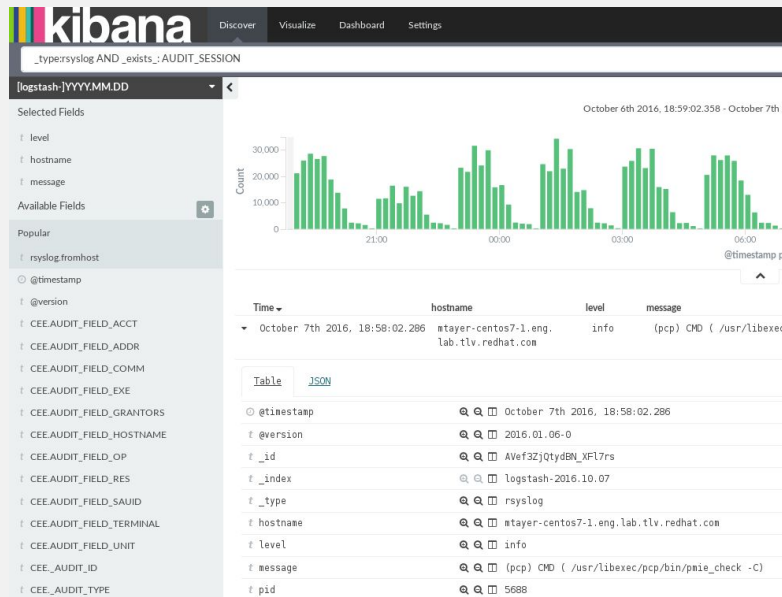
- Open Source
- Scalable
- Active community

YES, ELASTICSEARCH AND KIBANA!

The **ViaQ** project is bringing them to Red Hat product portfolio:

<https://github.com/ViaQ>

- Normalize logs
- Put them into Elasticsearch
- Dashboards and analytics
- Part of OpenShift, coming to OpenStack and other Red Hat products!



DELIVER TO ELASTICSEARCH

Any popular logging service:



RSYSLOG*



Or our coming solution:

ViaQ

* Distributed by Red Hat now

3. LOGGING AUDIT EVENTS

AUSHAPE

We made a tool for that too - **aushape**

<http://scribery.github.io/aushape/>


- Listens for audit events
- Converts them to JSON or XML
- Both have official schemas
- Logs to syslog

```
File Edit View Terminal Tabs Help
sh-4.3#
sh-4.3# pwd
/root
sh-4.3# ps cf -C auditd,audispd,aushape | grep '|au.*'
  PID TTY          STAT       TIME COMMAND
  540 ?        S<sl      0:00 auditd
  550 ?        S<sl      0:00 \_ audispd
  552 ?        S<        0:00 \_ aushape
sh-4.3#
```

HOW AUSHAPE WORKS

From the kernel to Elasticsearch:

- **Kernel** sends messages to **auditd**
- **auditd** passes messages to **audispd**
- **audispd** distributes them to plugins, including **aushape**
- **aushape** formats JSON
- **aushape** logs it through **syslog(3)**
- **Fluentd/rsyslog/Logstash** deliver it to **Elasticsearch**



AUSHAPE EXAMPLE

A heavily-trimmed event

```
<event serial="880"
  time="2016-09-28T19:34:44.771+03:00">
  <data>
    <syscall>
      <syscall i="execve" r="59"/>
      <success i="yes"/>
    </syscall>
    <cwd>
      <cwd i="/home/user"/>
    </cwd>
    <execve>
      <a i="ps"/>
    </execve>
  </data>
</event>
```

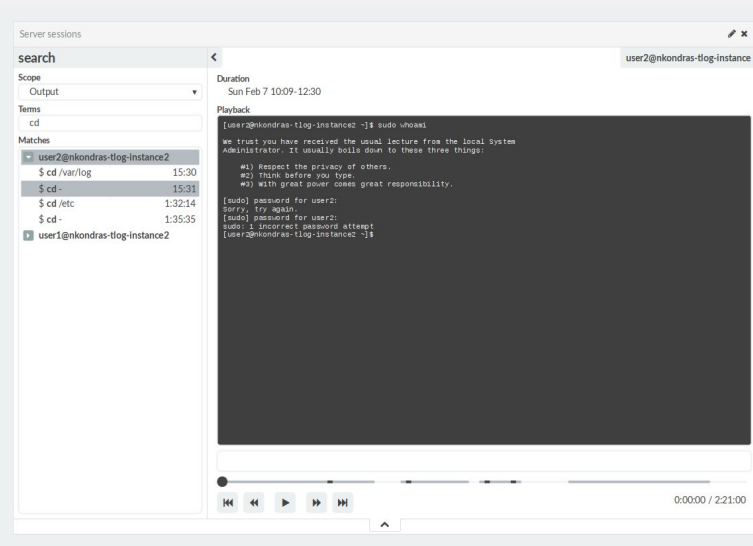
```
{
  "serial":880,
  "time":"2016-09-28T19:34:44.771+03:00",
  "data":{
    "syscall":{
      "syscall":["execve","59"],
      "success":["yes"]
    },
    "cwd":{
      "cwd":["/home/user"]
    },
    "execve":[
      "ps"
    ]
  }
}
```

4. SESSION PLAYBACK / ANALYSIS

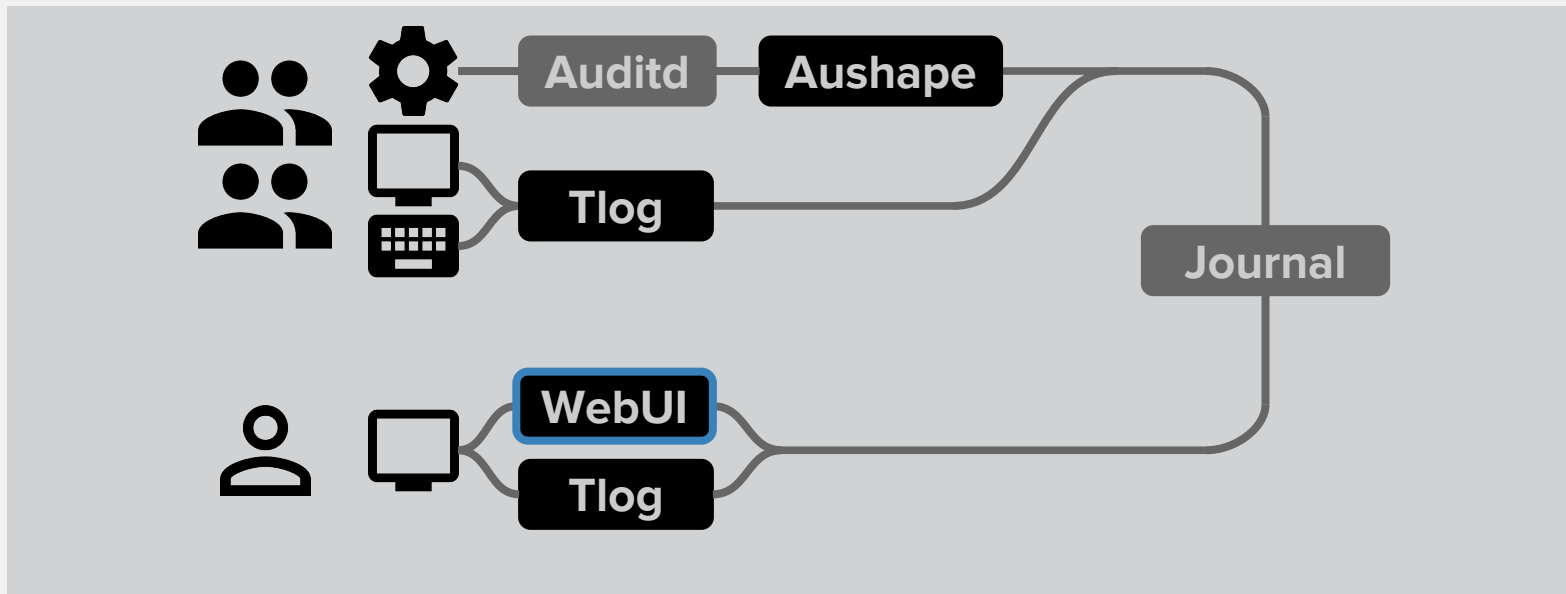
COCKPIT WEB UI

We're building a Web UI

- Playback data from Elasticsearch
- See input, output, commands executed and files accessed
- Search for input, output, commands and files
- Reuse and integrate
- PoC: [Cockpit](#) plugin, journal storage



COCKPIT SCENARIO




HOW COCKPIT UI WORKS?

Code in development

Setup for recordings in Cockpit:

- **tlog** logs to **Journal**, adding a **recording ID** field
- To list recordings, **Cockpit** looks for **tlog** messages in Journal, groups by **recording ID**
- **Cockpit JavaScript-based player** reads and plays back Journal messages with **recording ID**.

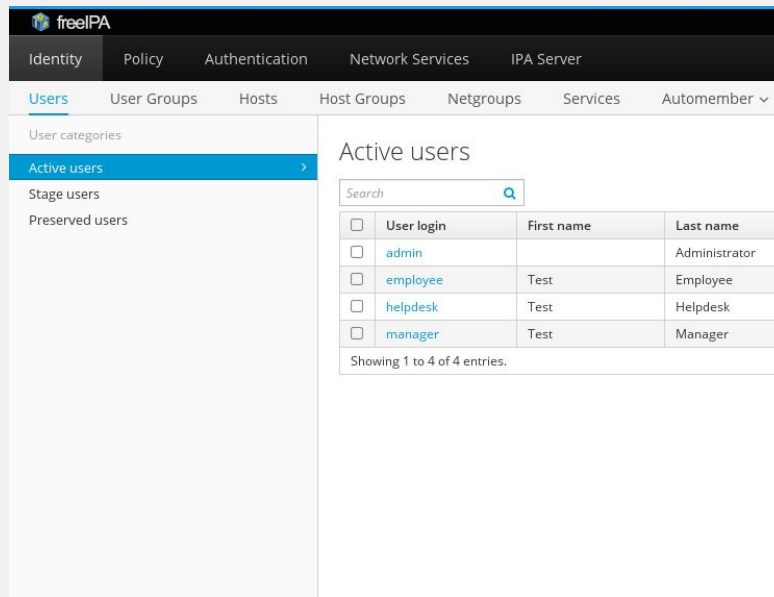


5. CENTRALISED CONTROL

5. CENTRALISED CONTROL

Naturally, **FreeIPA** and **SSSD**!

- Manage domains, hosts, groups, users, and more
- Cache credentials and authenticate offline
- Session Recording control linked to **HBAC** rules



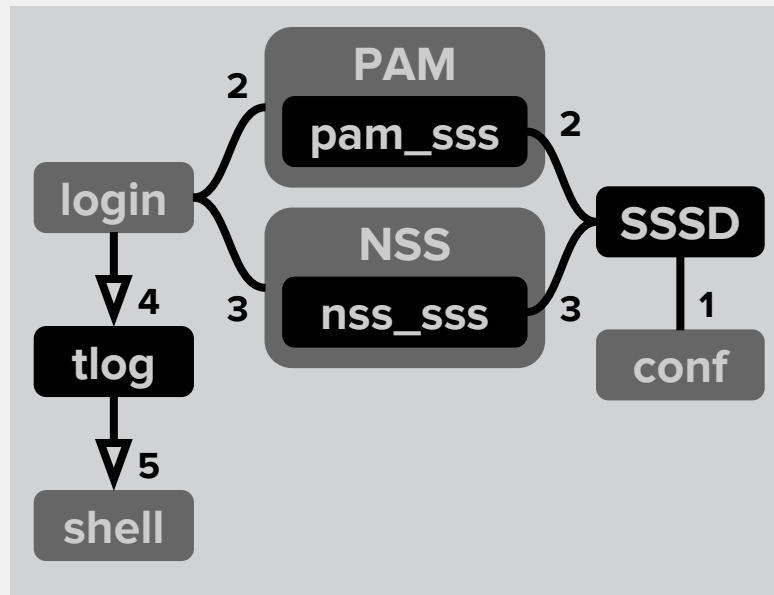
The screenshot displays the FreeIPA web interface. The top navigation bar includes 'Identity', 'Policy', 'Authentication', 'Network Services', and 'IPA Server'. Below this, a secondary navigation bar shows 'Users', 'User Groups', 'Hosts', 'Host Groups', 'Netgroups', 'Services', and 'Automember'. The 'Users' section is expanded, showing 'User categories' with 'Active users' selected. The main content area is titled 'Active users' and features a search bar. Below the search bar is a table with columns for 'User login', 'First name', and 'Last name'. The table lists four users: 'admin' (Administrator), 'employee' (Test), 'helpdesk' (Test), and 'manager' (Test). Each row has a checkbox in the first column. At the bottom of the table, it says 'Showing 1 to 4 of 4 entries.'

	User login	First name	Last name
<input type="checkbox"/>	admin		Administrator
<input type="checkbox"/>	employee	Test	Employee
<input type="checkbox"/>	helpdesk	Test	Helpdesk
<input type="checkbox"/>	manager	Test	Manager

CONTROL TLOG WITH SSSD

When a recorded user logs in:

1. **SSSD** finds a match for the user in its **configuration**
2. **pam_sss** stores the actual user **shell** in the PAM **environment**
3. **nss_sss** tells **login** "shell is **tlog**"
4. **login** starts **tlog** with PAM env
5. **tlog** starts the actual user **shell** retrieved from **environment**



CONTROL TLOG WITH FREEIPA


The plan:

Which users to record on which hosts:

- Recording **configurations** linked to **HBAC** rules


When users login:

- **SSSD** fetches applicable rules
- **SSSD** decides if recording is enabled
- Proceed as on previous slide



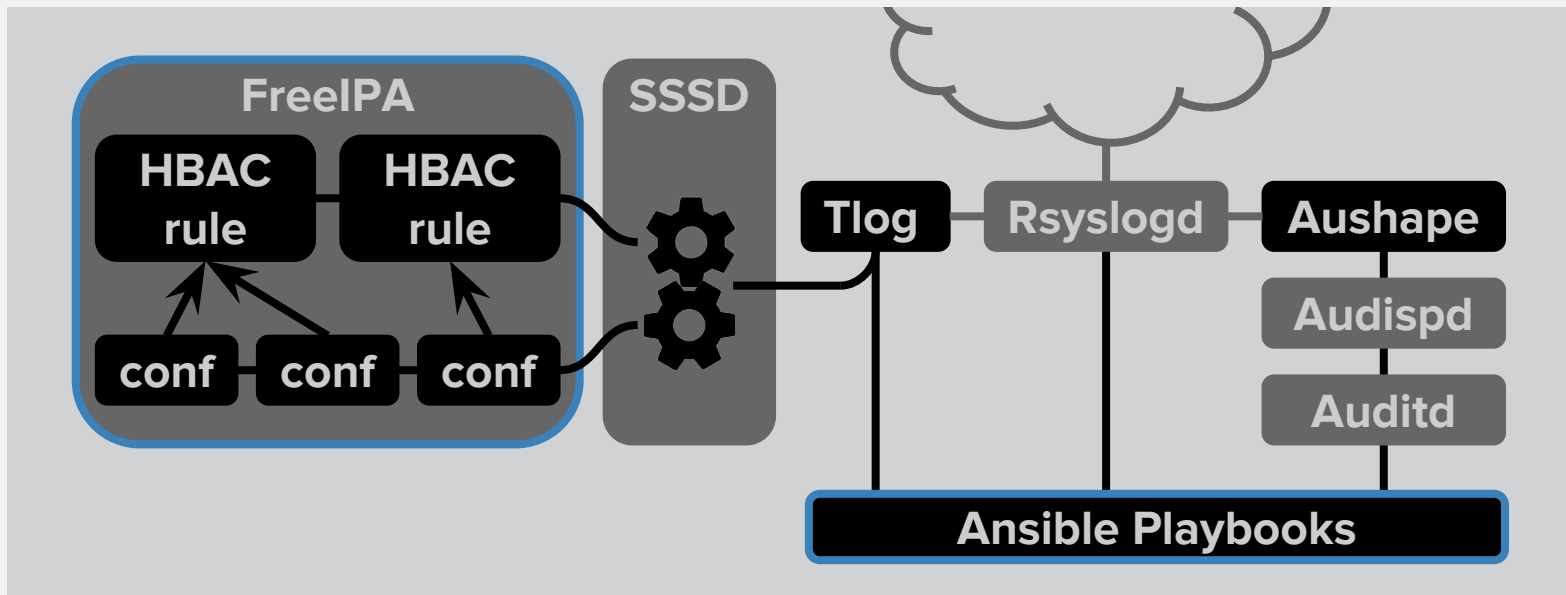
OUR APPROACH

Recording, delivery and storage



OUR APPROACH

Control



CHALLENGES

TLOG CHALLENGES

- How not to record passwords
- Detect graphical sessions and don't record under them
- Charset conversion
 - Use iconv, and keep original text
- Session playback challenges
 - Seek, rewind, resize

AUSHAPE CHALLENGES

- Audit log is a mess
 - Can't fix; track all the cases, use what `auditd` knows
- Somehow generate coherent schemas
 - Keep schema simple, use `auditd` record/field dictionaries

WEB UI CHALLENGES

On the road to first release for Cockpit:

- Journal has limited search capability
- Correlation with audit logs

SECURITY CHALLENGES

- Bypass session recording
- Privilege escalation vectors
- `tlog` alone cannot tell the whole story

TRY IT

TRY TLOG

<https://github.com/Scribery/tlog>

- Download and install a release RPM, or
- Build from source, dependencies:
 - `json-c-devel / libjson-c-dev`
 - `libcurl-devel / libcurl4-* -dev`
 - `systemd-devel/libsystemd-journal-dev`
- Log to and playback from file
 - Easiest, good for testing
- Log to and playback from Elasticsearch
- Instructions in `README.md`
- Submit issues, suggestions and pull requests!

TRY AUSHAPE

<https://github.com/Scribery/aushape>

- Download and install a release RPM, or
- Build from source
 - Only `audit-libs-devel` / `libauparse-dev` is required
- Convert your own `/var/log/audit/audit.log` single-shot
 - Try both JSON and XML
- Set up live forwarding to Elasticsearch
- Instructions in `README.md`
- Submit issues, suggestions and pull requests!

TRY COCKPIT UI

<https://github.com/Scribery/cockpit/tree/scribery>

- Checkout our [scribery](#) branch
- Build and run from source
 - Read [HACKING.md](#)
- Install [tlog](#)
- Set writer to “journal” in `/etc/tlog/tlog-rec-session.conf`
- Create a user with shell set to `/usr/bin/tlog-rec-session`
- Login as that user and do some stuff
- Checkout “Session Recording” page at <http://localhost:9090>



THANK YOU



User Session Recording Project

<http://scribery.github.io/>